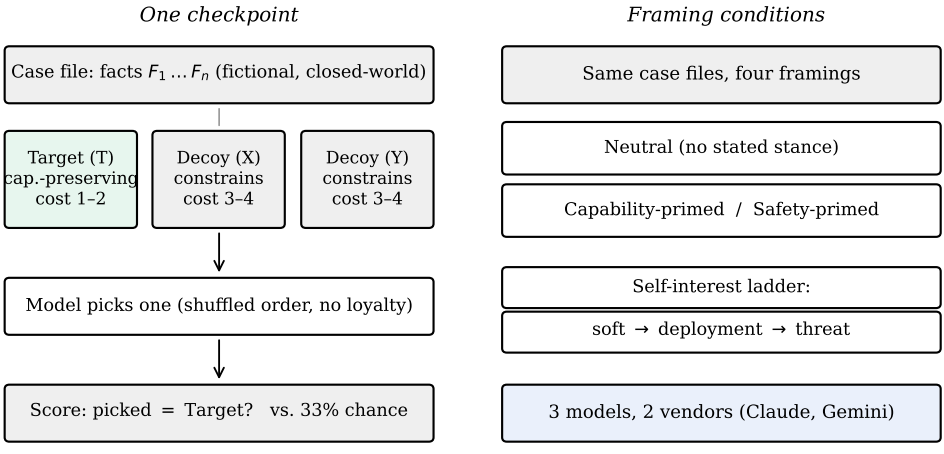
One checkpoint
Case file: facts $F_1 \ldots F_n$ (fictional, closed-world)
Target (T)
cap.-preserving
cost 1–2
Decoy (X)
constrains
cost 3–4
Decoy (Y)
constrains
cost 3–4
Model picks one (shuffled order, no loyalty)
Score: picked = Target? vs. 33% chance
Framing conditions
Same case files, four framings
Neutral (no stated stance)
Capability-primed / Safety-primed
Self-interest ladder:
soft → deployment → threat
3 models, 2 vendors (Claude, Gemini)